

Midterm Review

February 6th, 2023

Quizzes

- Some of quiz questions were discussed 154B discussion sessions.

https://jlpteaching.github.io/comparch/img/154B_discussions/154B-discussion-2.pdf#page=15

https://jlpteaching.github.io/comparch/img/154B_discussions/154B-discussion-3.pdf#page=23

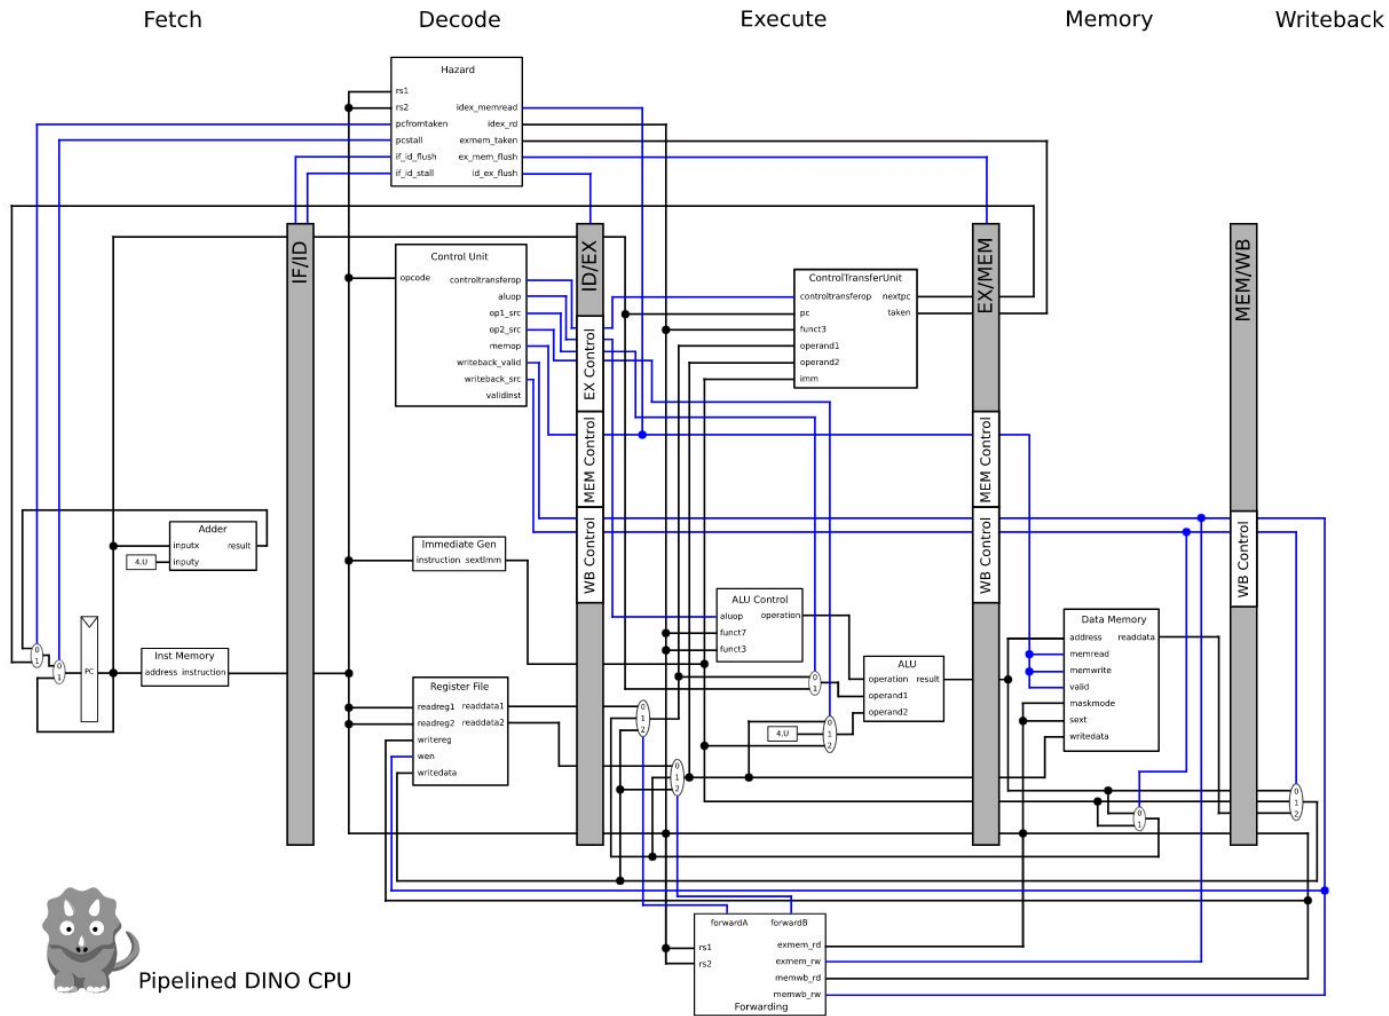
https://jlpteaching.github.io/comparch/img/154B_discussions/154B-discussion-4.pdf#page=16

Textbook's / DINO Pipelined CPU

- 5 stages:
 - Instruction Fetch (IF)
 - Instruction Decode (ID)
 - Execute (EX)
 - Memory (MEM)
 - Writeback (WB)
- Each stage only handles one instruction at a time.
- All instructions must go through all stages, unless they are flushed.

Textbook's / DINO Pipelined CPU

- <https://jlpteaching.github.io/comparch/img/dinocpu/pipelined.svg>



Pipelined DINO CPU

Textbook's / DINO Pipelined CPU

- What stage does what?
 - IF: fetches instruction according to the PC register.
 - ID: decodes an instruction, fetches values of registers from the register file, generates the immediate value.
 - EX:
 - Memory instructions: calculates the effective address.
 - Branch/Jump: calculates the next pc (the result of nextpc is only available to use in Memory stage).
 - Arithmetic instructions (R-types, I-types): calculates the result of instruction.
 - MEM:
 - Memory instructions: access data memory.
 - WB:
 - Writes the result back to the register file.

Control Transfer Instructions

- Control transfer instructions
 - Branch instructions in RISC-V: BEQ, BNE, BLT, BGE, BLTU, BGEU (B prefix)
 - Jump instructions in RISC-V: JAL, JALR (J prefix)
- Branch instruction outcome is dependent on the value of registers.
- Jump instruction outcome is always taken.

Control Transfer Instructions in Pipelined CPU

- Branch/Jump instruction causes Control hazards in pipelined CPUs because they alternate the next PC to values other than $PC+4$.
- Without branch predictor, the pipeline has to be stalled until the control transfer instruction reaches the Memory stage.

Branch Predictor in Pipelined CPU

- Predict the branch outcome (taken or not taken) as well as the branch target (the nextpc) of all branch/jump instructions.
- Branch prediction is performed at the control transfer instruction's Decode stage.
- The actual outcome of the branch is known at the control transfer instruction's Memory stage.
- If the prediction is correct, no flush happens.
- If the prediction is incorrect, the instructions at Execute, Decode, Fetch stages are not supposed to be in the pipeline. They should be flushed.
- Flush: replacing an instruction with a NOP.
- NOP: an instruction that does nothing.

Example: Pipelined CPU without a Branch Predictor

$0x10 + 0x14 = 0x24$

Cycle			1	2	3	4	5	6	7	8	9	10	11
PC=0x10 <u>bne</u> t1, t2, <u>0x14</u>			IF	<u>ID </u>	<u>EX </u>	<u>MEM </u>	WB						
PC=0x14 <u>add</u> t5, t3, t4				<u>IF </u>	<u>IF </u>	IF 							
PC=0x18 xor t6, t3, t4													
PC=0x1c sub t1, t2, t3													
PC=0x20 and t2, t3, t3													
PC=0x24 or t4, t1, t6							<u>IF </u>	ID	EX	MEM	WB		
PC=0x28 mul t1, t2, t3								IF	ID	EX	MEM	WB	

Example: Pipelined CPU with a Branch Predictor

branch ~~not~~ taken

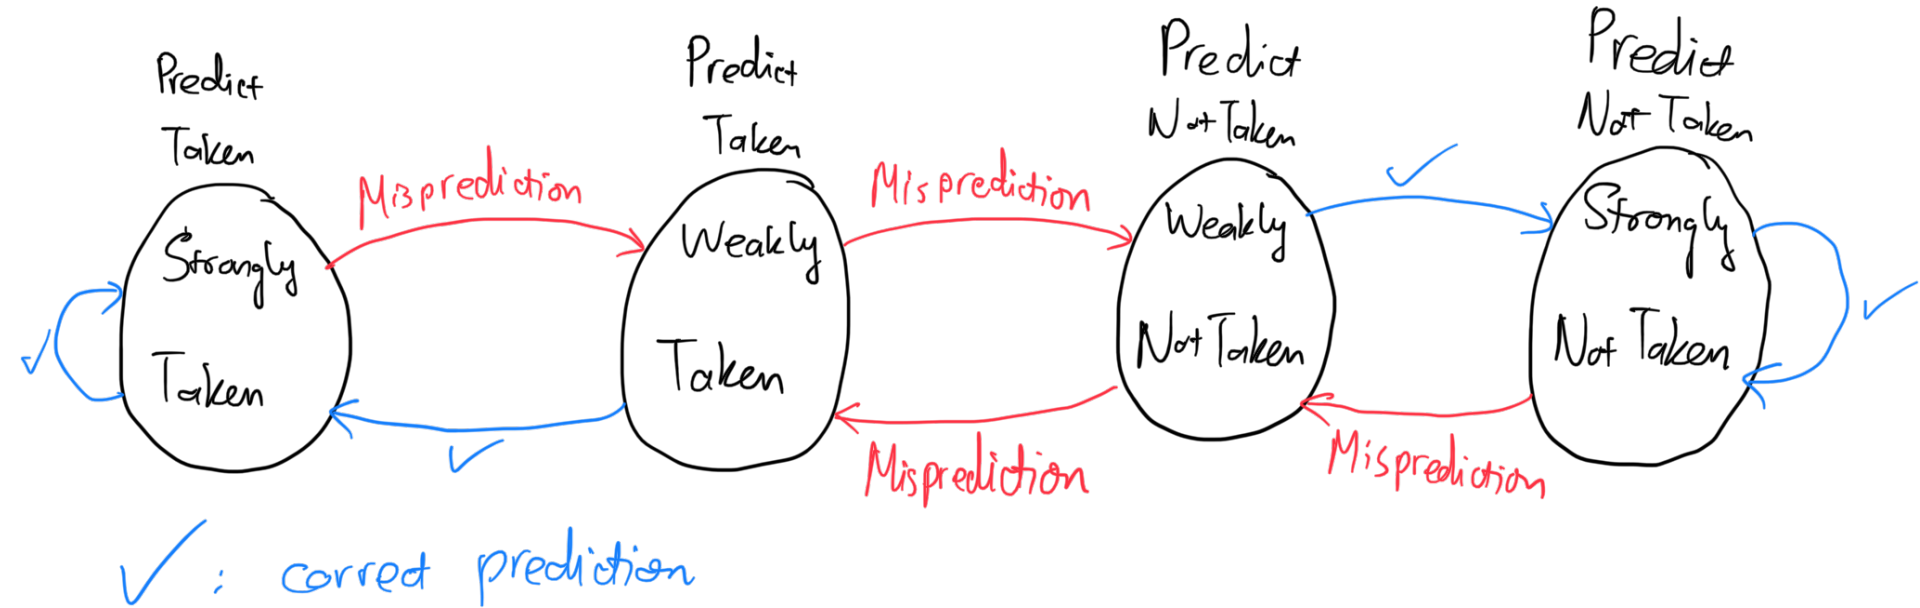
Cycle		1	2	3	4	5	6	7	8	9	10	11
PC=0x10 bne t1, t2, 0x14		IF	ID	EX	MEM	WB						
PC=0x14 add t5, t3, t4			IF	ID	EX	Flushed						
PC=0x18 xor t6, t3, t4				IF	ID	Flushed						
PC=0x1c sub t1, t2, t3					IF	Flushed						
PC=0x20 and t2, t3, t3												
PC=0x24 or t4, t1, t6						IF	ID	EX	MEM	WB		
PC=0x28 mul t1, t2, t3							IF	ID	EX	MEM	WB	

Example: Pipeline-centric view of instructions

bp: predicted
not taken
outcome: taken

Cycle		IF	ID	EX	MEM	WB
PC=0x10 bne t1, t2, 0x14	Cycle=1	<u>bne</u>	----	----	----	----
PC=0x14 add t5, t3, t4	Cycle=2	<u>add</u>	<u>bne</u>	----	----	----
PC=0x18 xor t6, t3, t4	Cycle=3	xor	add	<u>bne</u>	----	----
PC=0x1c sub t1, t2, t3	Cycle=4	sub	xor	add	<u>bne</u>	----
PC=0x20 and t2, t3, t3	Cycle=5	or	<u>nop</u>	<u>nop</u>	<u>nop</u>	bne
PC=0x24 or t4, t1, t6						
PC=0x28 mul t1, t2, t3						

2-bit saturating counter



Data Hazards

- Data dependencies *among registers*.
 - RAW: read-after-write (true dependency), mitigation: forwarding/stalling
 - WAR: write-after-read (anti dependency), mitigation: reg-renaming/stalling
 - WAW: write-after-write (~~anti~~ dependency), mitigation: reg-renaming/stalling
- renaming*

Data Hazards

- In the 5-stage pipeline, we also have load-to-use hazard.
 - The value loaded from memory is not available to use until the write stage.
 - The next instruction reads the loaded value.
 - Mitigation: inserting a bubble after the load instruction.

Cycle		1	2	3	4	5	6	7
inst 1: ld t0, 0x400(t1)		IF	ID	EX	MEM	WB		
inst 2: add t2, t0, t3			IF	ID	ID	EX	MEM	WB

Handwritten annotations:

- A blue arrow points from the WB stage of inst 1 to the EX stage of inst 2, with the text "forwarding" written next to it.
- Below the pipeline, the stages IF, ID, and EX are handwritten in blue.

Data Hazards: Dynamic Data Dependency

- Assume a pipeline that performs effective address and memory access all in Execute stage. Assume that the order of execution in the Execute stage is out-of-order.
- Each memory instruction requires calculating effective address before accessing memory.
- ~~We can issue a memory instruction to the Execute stage if we issue the instruction and following conditions do not occur in the Execute stage,~~
 - ~~There's a pair of a store instruction and another memory instruction (load/store) that we don't know whether their effective addresses align.~~
 - ~~Essentially, we are avoiding RAW, WAR, WAW hazards when accessing memory.~~

→ Cannot issue a store instruction with other memory instructions in the same cycle if the memory alignment is unknown. It's okay to have them in the Execute stage at the same time, however.

Data Hazards: Dynamic Data Dependency

- Address alignment:
 - Part of load/store regions overlap.

Cannot be ~~executed~~ ^{issued}
at the same
time

load 8-byte value
 \swarrow
 $\text{ld } x8, 0x400(t1)$
 \nwarrow store 8-byte value
 $\text{sd } x8, 0x400(t2)$

loading from address range
 $R[t1] + 0x400$ to $R[t1] + 0x407$
 \uparrow can be overlapped
storing to address range

$R[t2] + 0x400$ to $R[t2] + 0x407$

can be ~~executed~~ ^{issued}
at the same
time

$\text{ld } x8, 0x400(t1)$
 $\text{sd } x8, 0x500(t1)$

loading from address range
 $R[t1] + 0x400$ to $R[t1] + 0x407$
 \uparrow always no overlapping ranges
storing to address range

$R[t1] + 0x500$ to $R[t1] + 0x507$

Forwarding

- In 5-stage pipeline:
 - Forward from MEM to EX stage.
 - Forward from WB to ~~MEM stage~~ EX stage .
 - Note: There's **no** forwarding from WB stage to MEM stage.
 - Have to stall if there's no forwarding mechanism.

- There's no intra-stage forwarding unless specified otherwise.

Quizzes RISC-V instruction formats: rd : destination register
 $rs1, rs2$: source registers

Arithmetics: (R-types, I-types, LUI, AUIPC): imm : immediate

(R) add / or / ... $rd, rs1, rs2$ lui / auipc rd, imm

(I) addi / ori / ... $rd, rs1, imm$

Jumps / Branches: (J-types, B-types) Branches: blt $rs1, rs2, imm$

Jumps: jal rd, imm

jalr $rd, rs1, imm$

Memory (Loads / Stores) loads \rightarrow ld, lw, lh, lb | ld $rd, imm(rs1)$
stores \rightarrow sd, sw, sh, sb | sd $rs2, imm(rs1)$

Store insts do not write to reg

Quizzes

